

gen-i



SMART ENERGY

AWS, some small hints for Dynamo DB and Lambda

David Grgič

2023-06-15

gen-i

- Sun (Sonce)...
- Retail...
- Trading:
 - gen-i do not own production or storage assets.
 - Data driven business model.
 - Trading based on:
 - Technical Analysis (RSI, MACD, etc.)
 - Fundamental Models:
 - ...
 - **Market Simulator**
 - ...

gen-i Market Simulator

- Input:
 - Power consumption.
 - Power plants with their technical parameters, availability, etc.
 - Fuel and CO2 prices.
 - Connection between countries (price zones).
 - Flow Based Market Coupling data.
- Output:
 - Future prices for most European countries (price zones).
- Challenges:
 - MILP type of a problem with several thousands of optimisation variables per each simulated hour.
 - Simulating up to end of next year delivery periods, e.g. more than ten thousand hours.
 - Scenarios to estimate risk associated with several fundamental parameters: $\prod n_i$
 - Compute intensive task, where results are required ASAP, but simulations are run sporadically:

⇒ **AWS: Lambda, DynamoDB, Step Functions**

AWS DynamoDB, introduction

- Fully managed.
- NoSQL database.
- Document database, data is stored in a flexible JSON-like format.
- Less structured than traditional SQL databases.
- Fast and predictable performance.
- Seamless scalability.

DynamoDB read (write) capacity

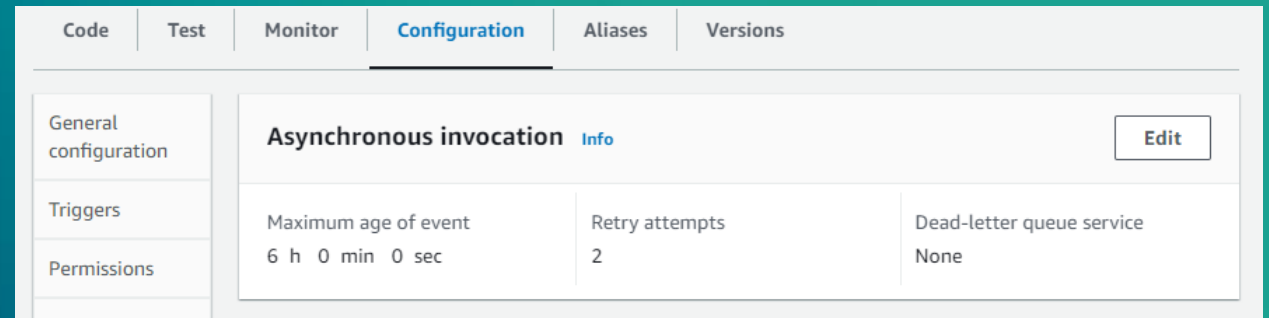
- Why it is important:
 - Related to costs!
 - If exceeded, operation (read or write) can be rejected.
- Reading data via „Query“ and „Scan“:
 - „ProjectionExpresion“ to limit retrieved attributes.
- Findings:
 - „ProjectionExpression“ does not decrease read capacity, since consumed read capacity is derived from total item size!
 - Read capacity is not derived from read items but from largest item in the database, even if that particular item is not part of read items!

AWS Lambda, introduction

- Serverless compute service:
 - Run your code without provisioning or managing servers.
- Can be event triggered: S3, HTML request, DynamoDB, etc.
- Supported variety of programming languages:
 - Python,
 - Java,
 - .NET,
 - etc.
- High availability.
- Automatic scaling.

Lambda provisioning

- Available lambdas:
 - Each AWS account has some number of Lambdas, available per AWS location:
 - This number of Lambdas can be increased by request, without costs.
- Retry for asynchronous invocations:
 - When throttled due to:
 - "Too Many Requests" or
 - "ThrottlingException"
 - Up to two times (1x initial + 2x retry)
 - Exponential backoff delay between retries.



Lambda provisioning

- Reserved concurrency:

- If function has defined reserved concurrency, than:
 - This function can use only up to reserved number of lambdas!
 - Remaining lambdas for other functions are decreased accordingly.
- Max: account lambdas - 100

- Burst concurrency (lambda scaling limitation):

- What does it mean:
 - Even if you have ten thousands of lambdas per account, they are not available instantly.
 - First minute you have, e.g. 1000* lambdas, each consequent minute you have 500 more lambdas: 1000; 1500; 2000; 2500; 3000; ...
 - If load decrease, e.g. number of used lambdas, you again have a limited ramp for increase (500 more each minute).
- * Initial burst concurrency is location dependent (3000: us-east-1, us-west-2, eu-west-1; 1000: **eu-central-1**, etc.; 500: other)
- Reserved concurrency does not save you from burst concurrency limitation!
- Can not be changed!

