

# **AWS SageMaker in Production**

Mitja Hofer





- 1. What is AWS SageMaker?**
- 2. Why SageMaker?**
- 3. Deploying models**
- 4. Bring your own model**
- 5. Deploying into production environment**
- 6. Local development**

# What is AWS SageMaker?

- A suite of tools for discovery, building, training and serving machine learning models
- SageMaker Studio
- Model and endpoint development and deployment
- Model serving endpoints

# Why do we use AWS SageMaker?

- The L3 step of AI initiative, creative scoring, requires a real-time inference model server
- An alternative to building our own Python environment with a http server, an inference model server, all supporting code, a testing and deployment pipeline.
- Support for GPU instances
- SageMaker allows for training of our model to be implemented there as well.
- SageMaker SDK

- Real-time inference
  - A managed rest API endpoint with support for payloads up-to 6MB and 60s processing time
- Serverless inference
  - For irregular traffic, 4MB payloads, 60s processing time
- Asynchronous inference
  - Queue based processing. Up to 1GB payloads and 1h processing time
- Batch transform
  - Offline processing of workloads - large datasets and days of processing time

# Bring your own model (or container)

- You may create your own model in a Python machine learning framework of your choice and supply it to SageMaker
- You need to implement your inference codebase
- inference.py implements a chain of python functions
- First you need to load your model
- Then the serving is broken down into three steps:
  - Input processing
  - Prediction
  - Output processing

- `model_fn(model_dir: str)`
  - `model_dir` is the `TargetModel` parameter of the `InvokeEndpoint` API call
  - It respects the `SAGEMAKER_SUBMIT_DIRECTORY` environment variable, that can point to local filesystem or S3 path

- `input_fn(request_body, request_content_type)`
  - Takes in a HTTP request and deserializes it into an object
- `predict_fn(input_object, model)`
  - gets the object and performs inference against the loaded model
- `output_fn(prediction, response_content_type)`
  - takes the prediction and forms a HTTP response



# Single vs Multi Model endpoints

- Real-time inference supports single and multi-model endpoints
- .tar.gz models
- Single model endpoint may still load multiple model files if you implement `inference.py` that way
- Multi model endpoint will dynamically load and offload models based on usage - Multi Model Server
- The user specifies the S3 location of the model files
- `InvokeEndpoint` API, `TargetModel` parameter
  - Allows uploading new model files without even touching SageMaker
- Similarly-sized models work best

- Three constructs of a SageMaker model deployment
  - **Model**
    - Specifies the inference script and model files location
    - Specifies the role that SageMaker runner assumes
    - Environment variables
  - **EndpointConfiguration**
    - Specifies the **Model**
    - Instance type, autoscaling configuration, routing strategy
  - **Endpoint**
    - Specifies the **DeploymentStrategy**
    - Serves your requests

# Deploying to Production Environment

- SageMaker SDK for Python for simple model development and deployment
- The `Model()` constructor takes the endpoint name parameter - a hard coupling between the model and the endpoint
- No option to update endpoint
- You are expected to have model files on S3, but SageMaker SDK reuploads the files with correct permissions to a new bucket

# Deploying to Production Environment

- SageMaker SDK hides away the three constructs and so you cannot update your endpoint, you must always recreate it
- What about AWS CDK?
- L1 vs L2 constructs
- Fits into infrastructure code
- Define Cloudwatch alarms for scaling and model monitoring

- Deploying inference code and model files using AWS CDK
- The instance that is part of the `Model` construct is long-lived
- Model files are loaded on demand
- Inference code is loaded on creation
- Two options to upload files to S3 using CDK:
- `S3.Asset`
  - Single file, named after the content digest
- `BucketDeployment`
  - Respects directory structure, may keep old files on S3

**What about local development?**

Celtra

# What about local development?

- Did you know that you can simply redirect any AWS SDK client to any domain you want?

```
new SageMakerRuntimeClient({  
    endpoint: 'foo.test'  
})
```

- This will make all requests from the SDK go to `foo.test`
- Local development?

# What about local development?

- SageMaker `Model` is just a managed instance for which you specify the container image it runs
- Add the same image that SageMaker uses in the cloud to your `docker-compose.yml`
- Bind mount `inference.py` and model files
- Add a reverse proxy to correctly handle model registration and rewriting of requests (headers)
- As SageMaker uses Multi Model Server in the background, you need to understand how it works



# What we did

- Use AWS SageMaker for serving a in-house inference model
- Use a multi-model realtime inference endpoint
- Write our own inference code that loads the model, handles http requests and calls the model
- Use AWS CDK to create the three constructs needed for model serving, and upload to S3 in two different ways
- Support local development by using the same SageMaker container we are using in production

| Celtra