

Reducing cold start delays in FaaS and Serverless Environments



About me


- I'm **Teodor J. Podobnik** from Slovenia 🇸🇮
- **SRE** at Prewave
- **Researcher** at Faculty of Computer Science and Information Science, University of Ljubljana
- **Author** of [eBPFChirp](https://ebpfchirp.substack.com)



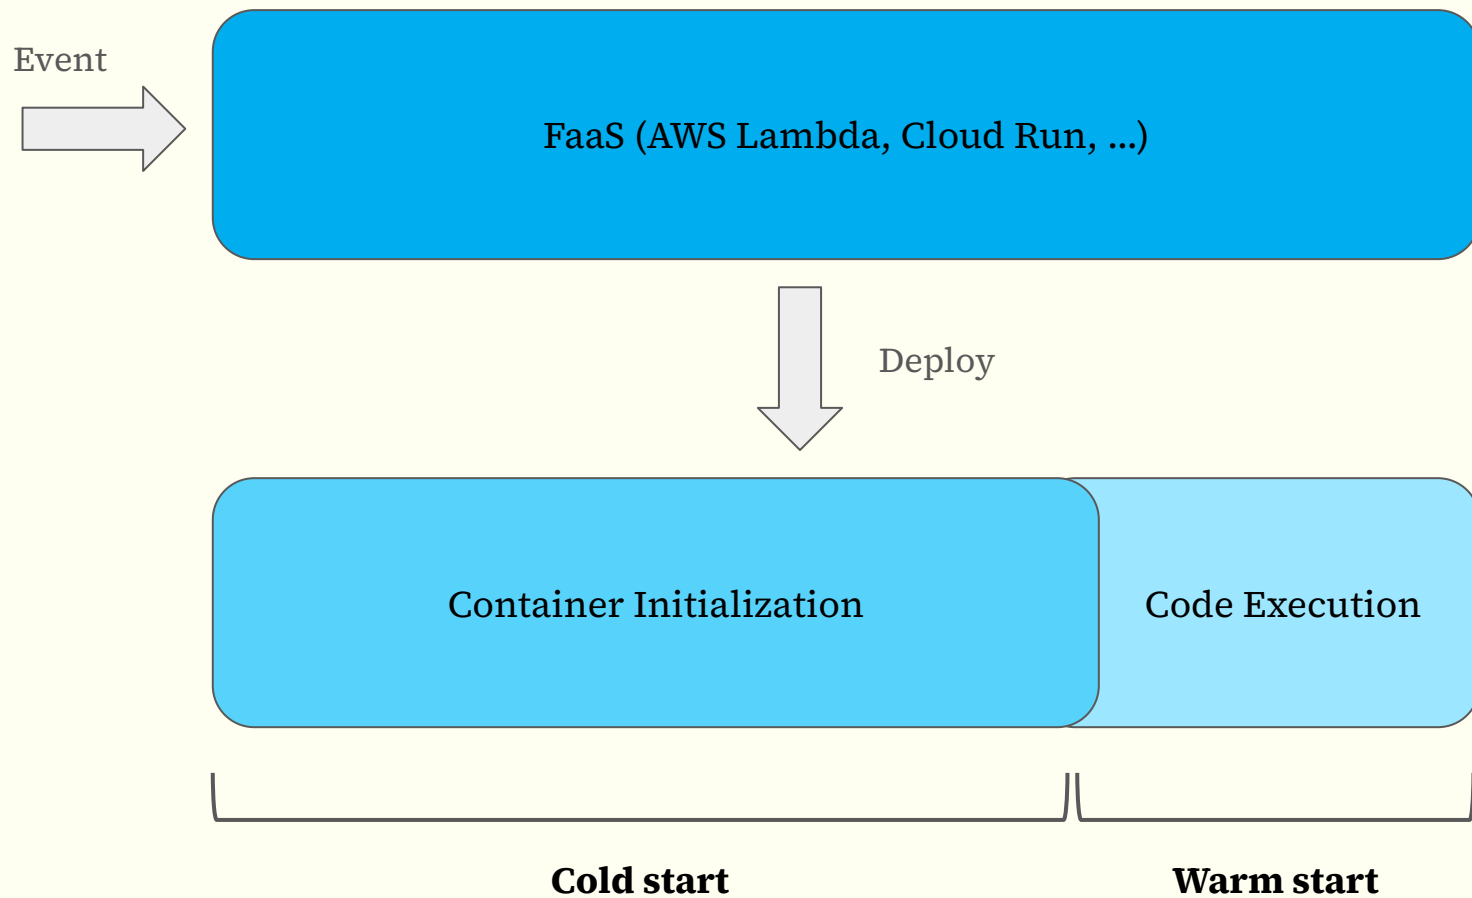
<https://www.prewave.com>
<https://fri.uni-lj.si>
<https://ebpfchirp.substack.com>

eBPF in Short

- Runs Directly in the Kernel
- Dynamically Attached at Runtime
- Minimal CPU Overhead
- Independent of User-Space Applications
- Use Cases:
 - Cloudflare: DDoS Protection
 - Meta: Load Balancing
 - Netflix: Net. Traffic Observability Enrichment



The Problem of cold start delays



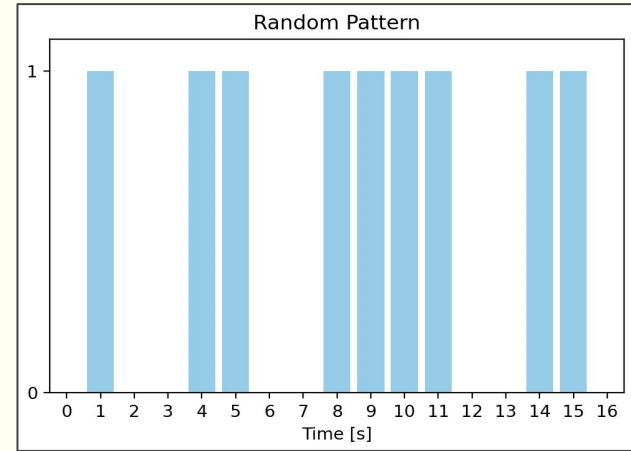
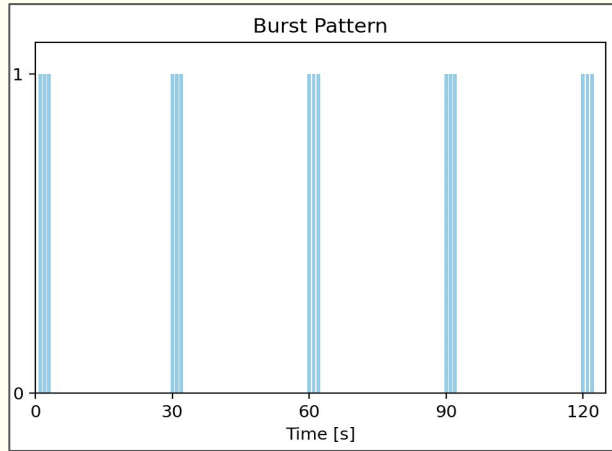

Cost vs. Latency

Real-time data analytics

Image or video transcoding

Web API endpoints

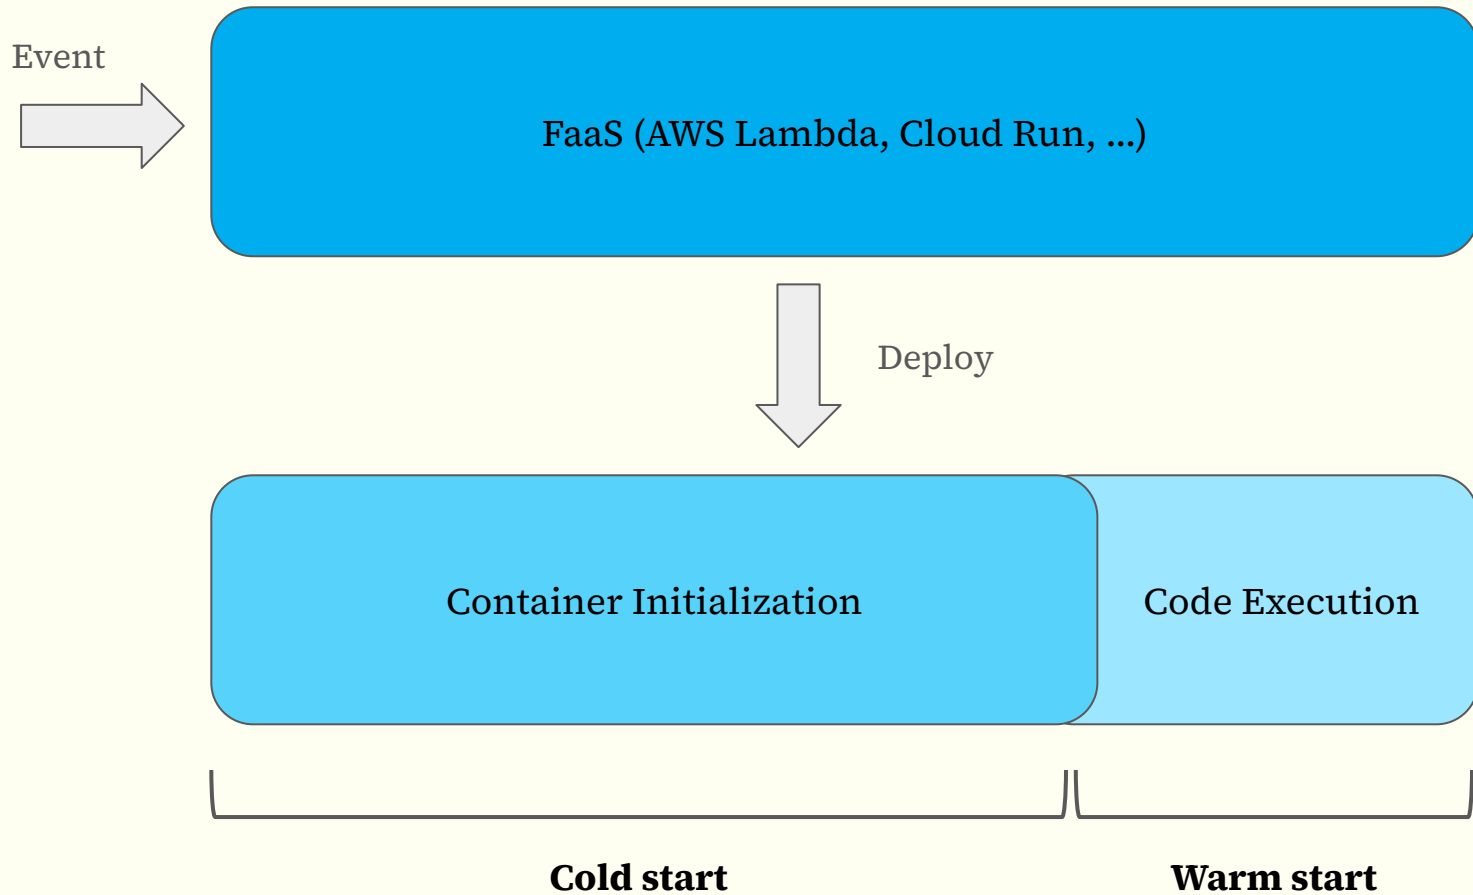
 **Demand**

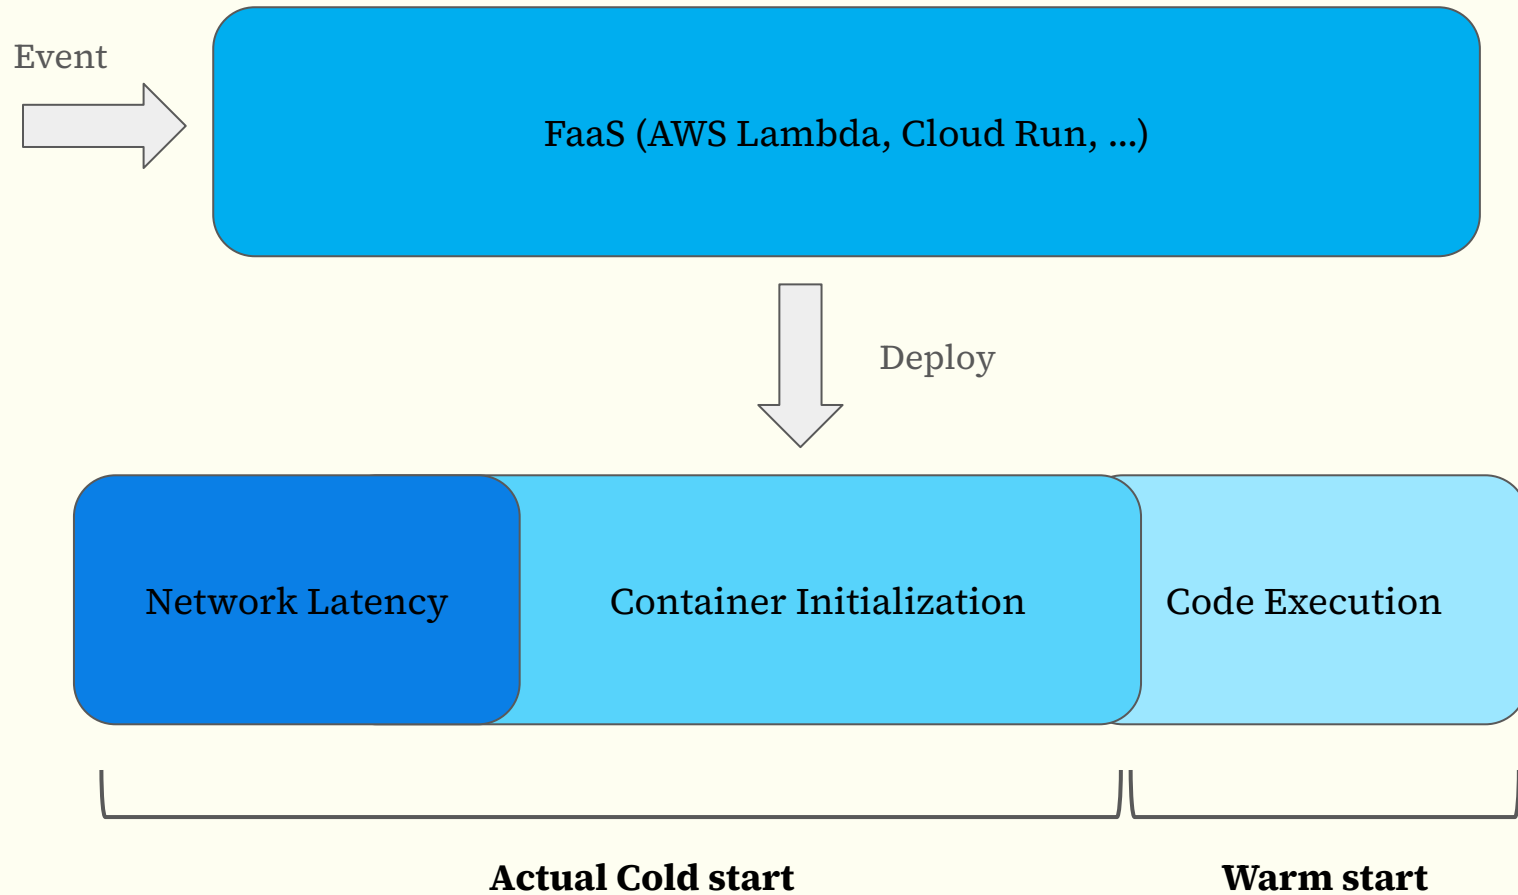


Constant demand,
no cold starts

Predictable demand,
no cold starts

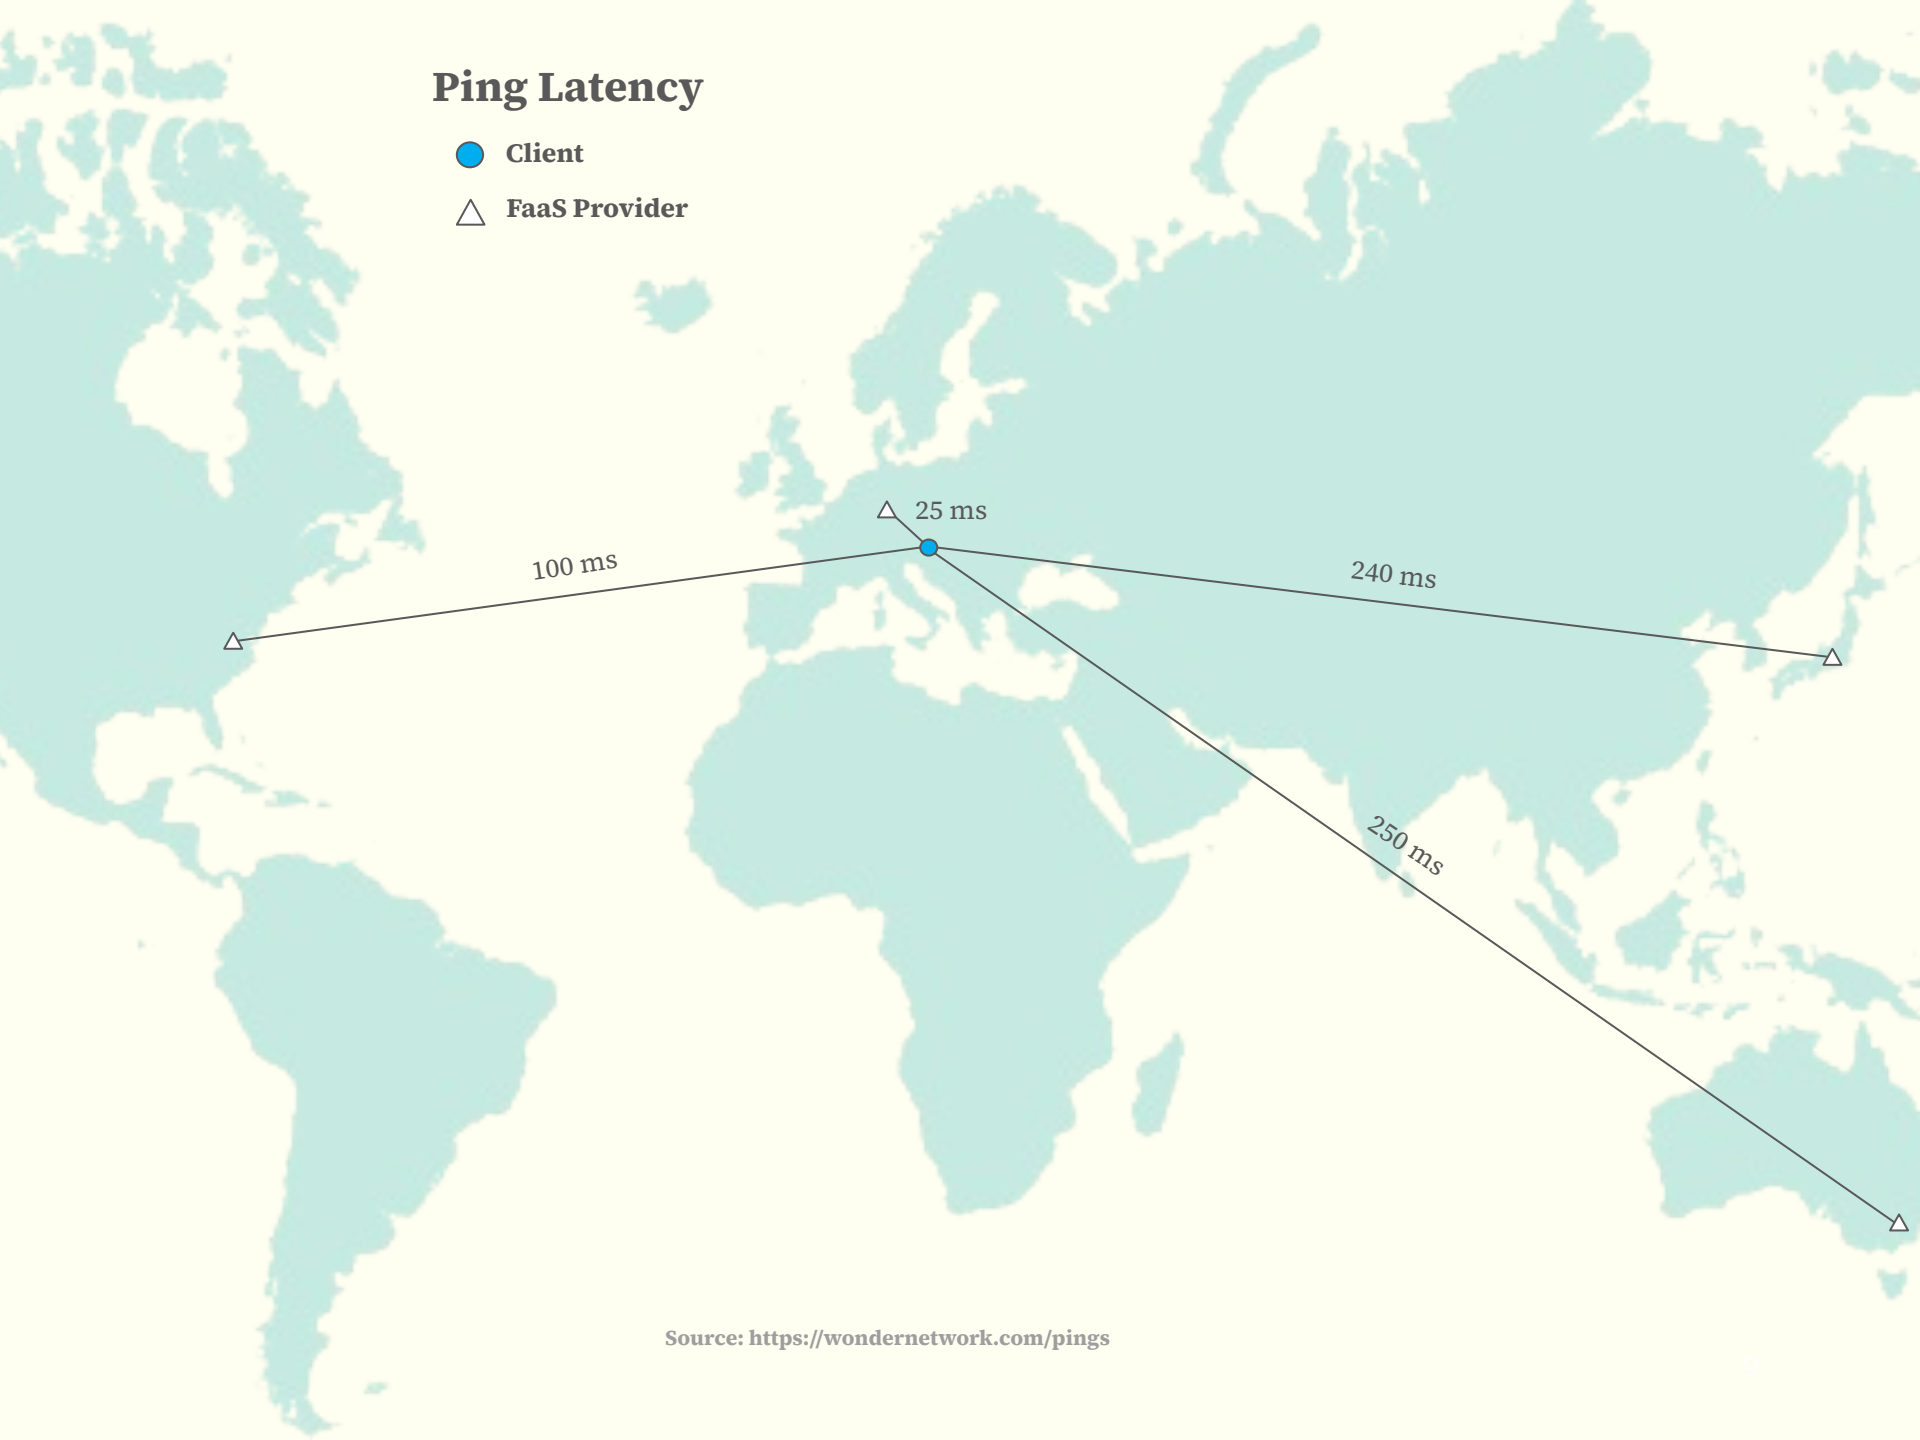






Ping Latency

- Client
- △ FaaS Provider

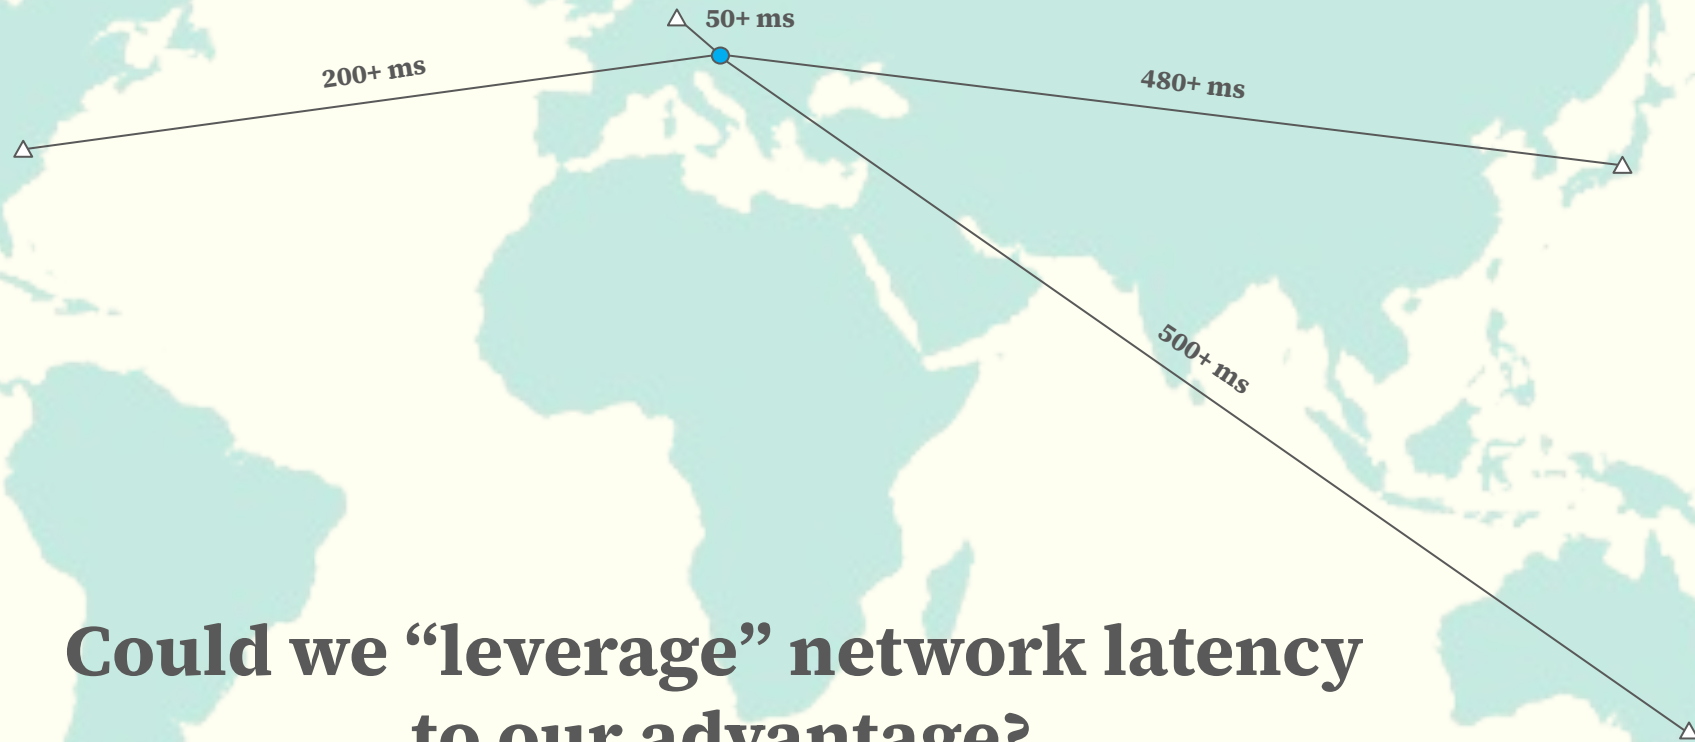


Source: <https://wondernetwork.com/pings>

Aprox. FaaS Latency

● Client

△ FaaS Provider

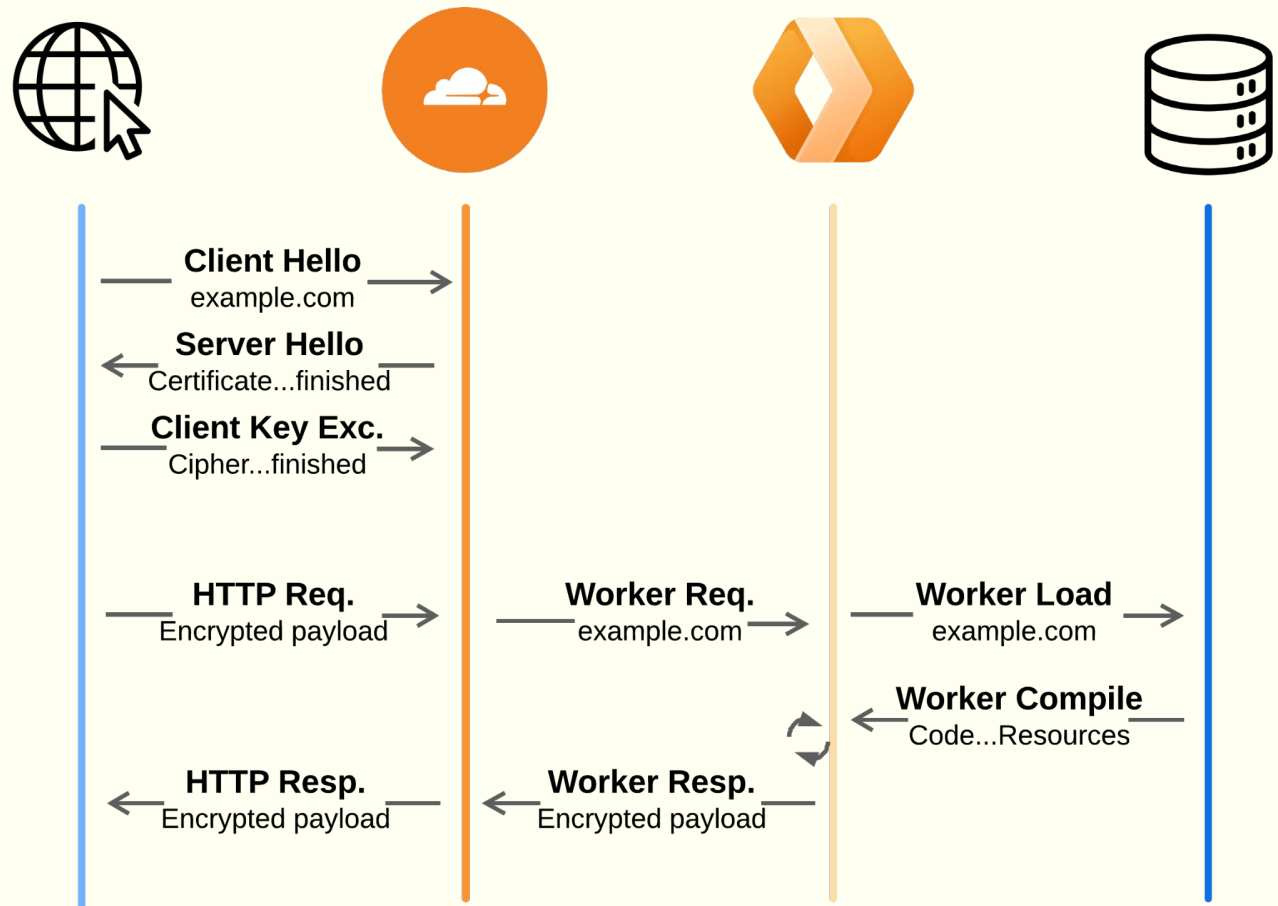


Could we “leverage” network latency to our advantage?

Source: <https://wondernetwork.com/pings>

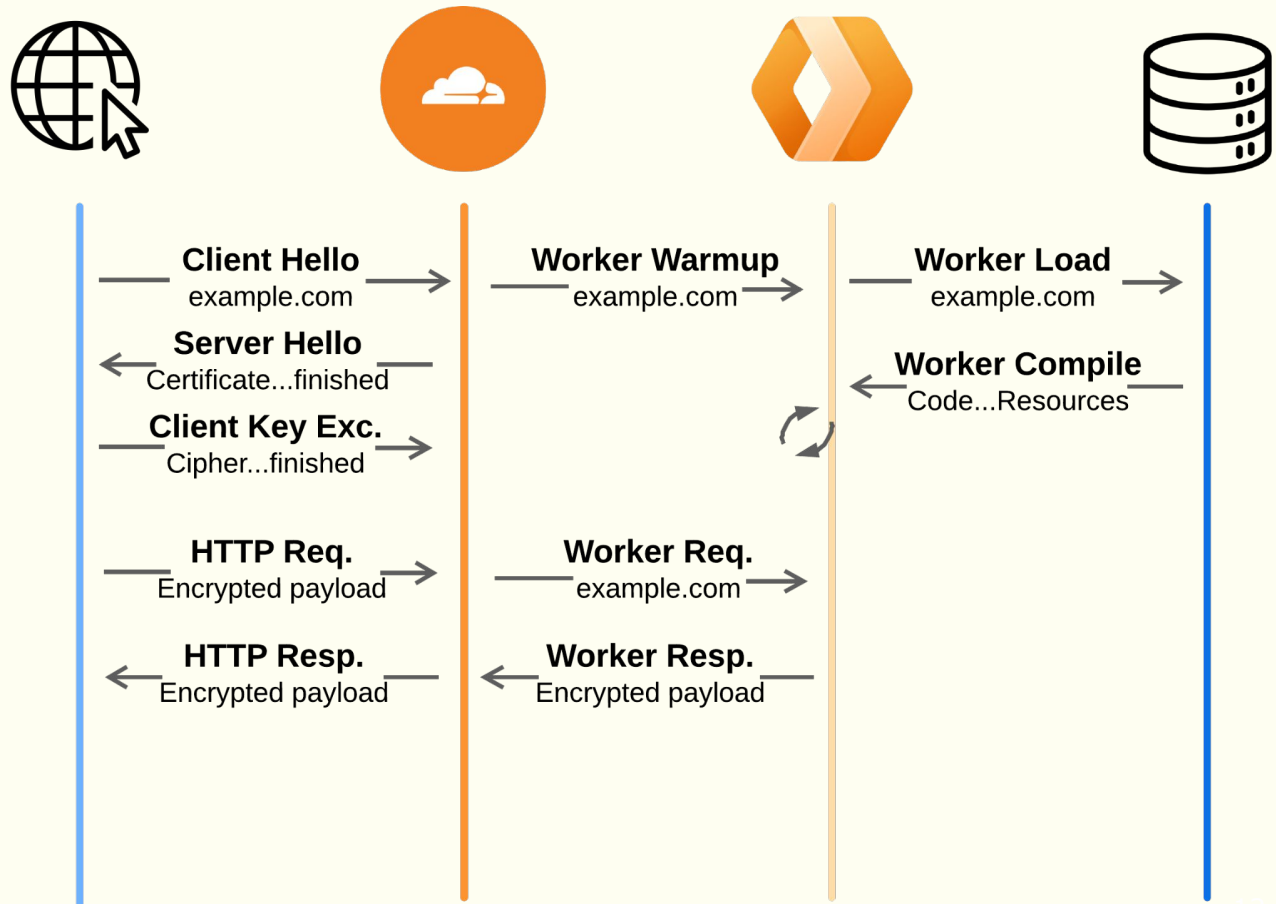
Cloudflare Solution

- “Pre-warming” during TLS handshake -> hostname from SNI
- CloudFlare Workers (NOT containers!)
- Not open-source



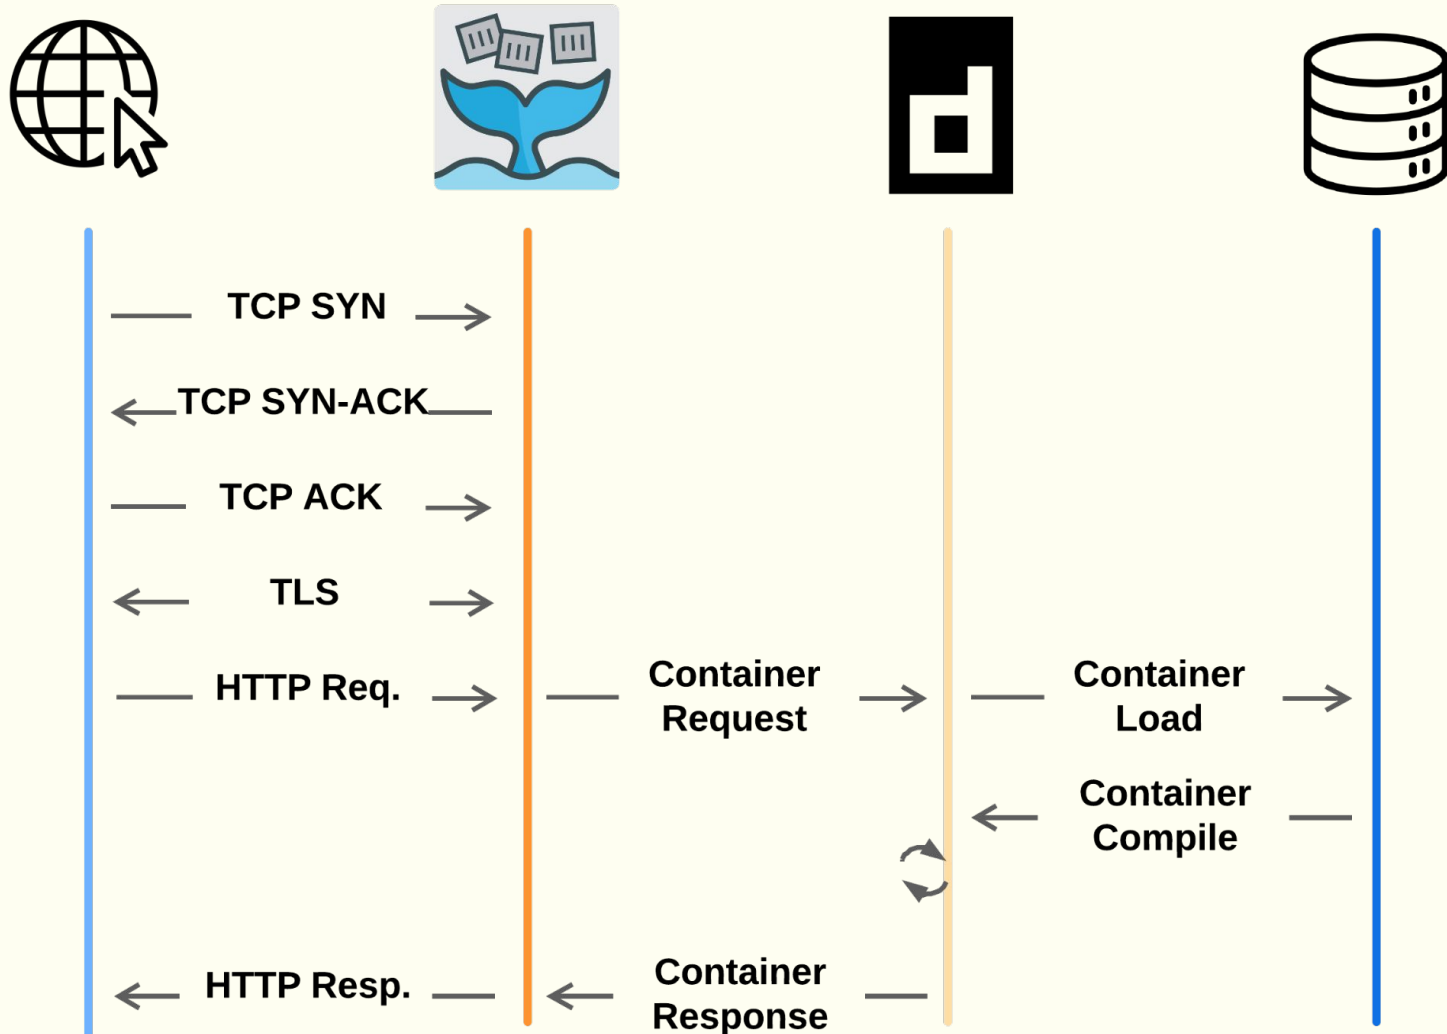
Cloudflare Solution

- “Pre-warming” during TLS handshake -> hostname from SNI
- CloudFlare Workers (NOT containers!)
- Not open-source

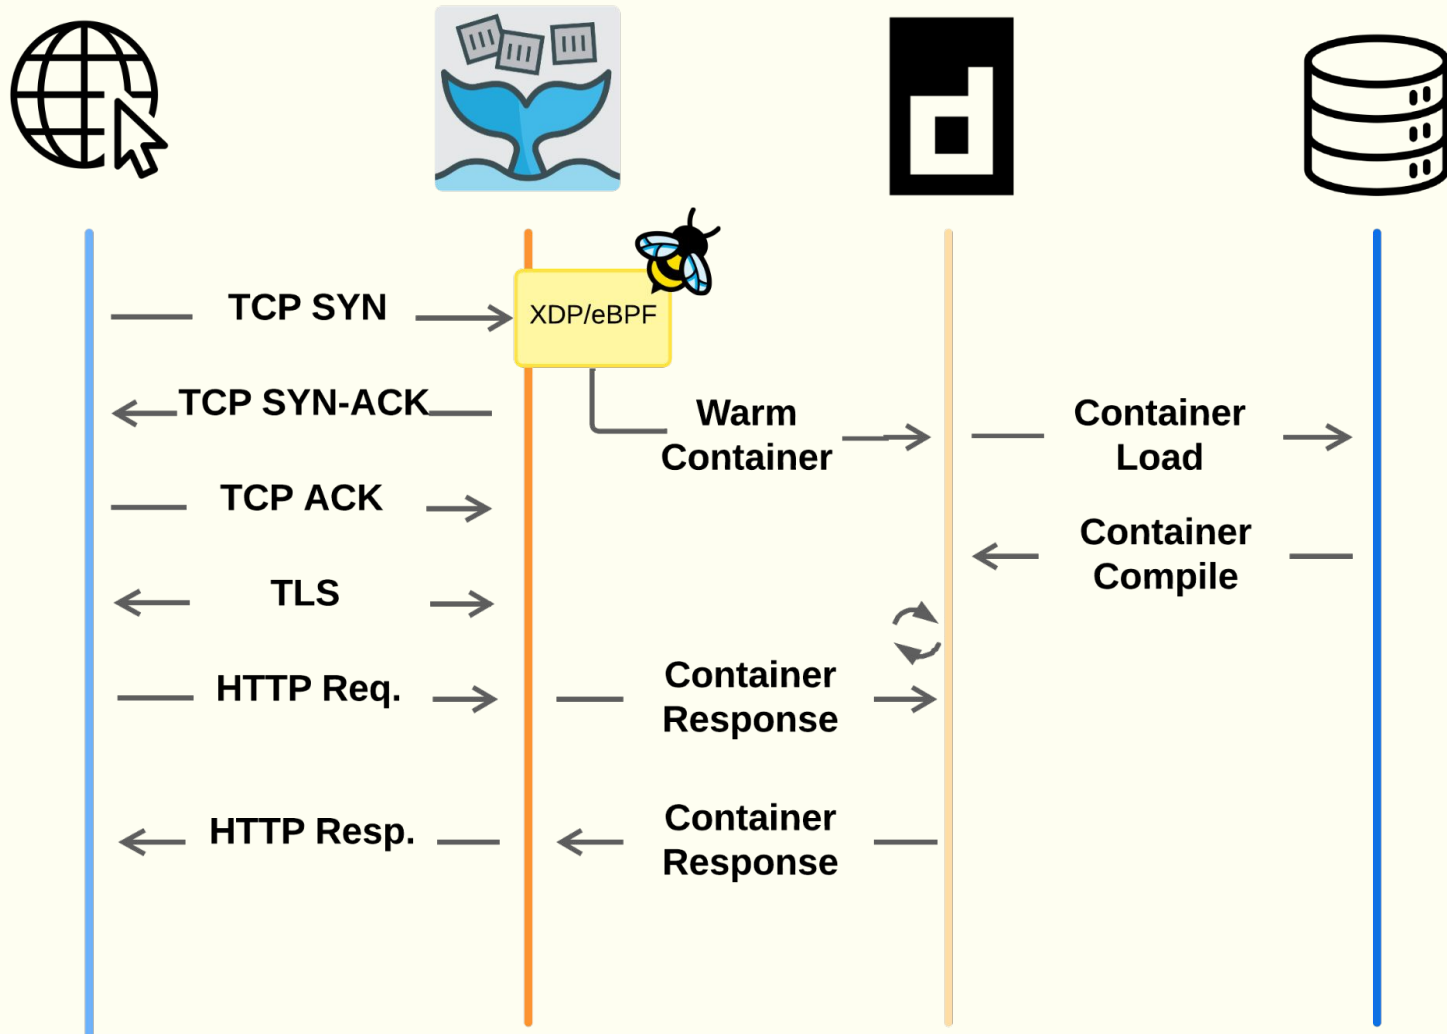


eBPF Solution

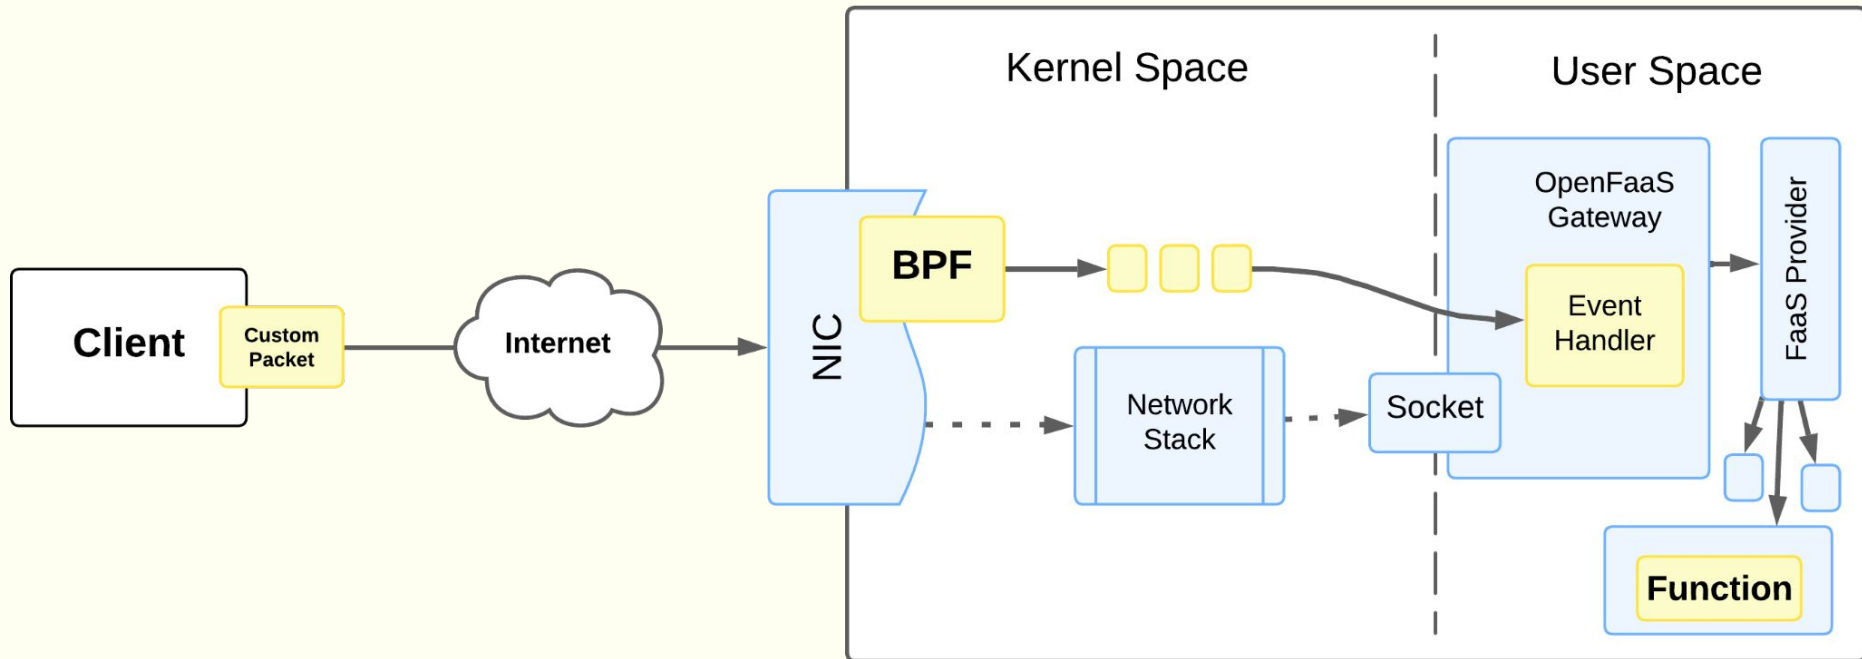
eBPF Solution



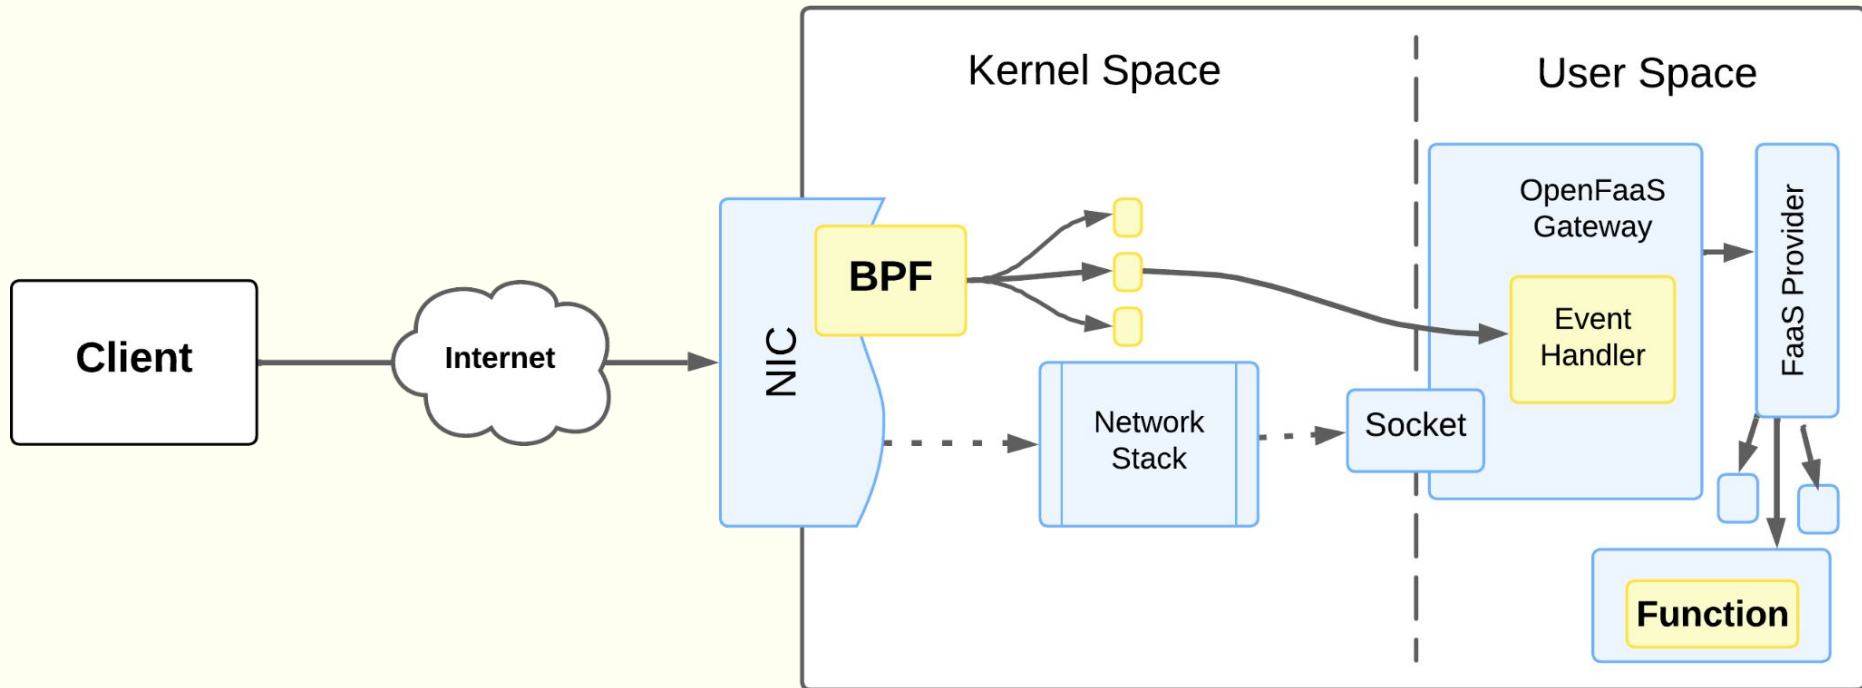
eBPF Solution



TCP Packet Custom Payload



FaaS Function per TCP Port

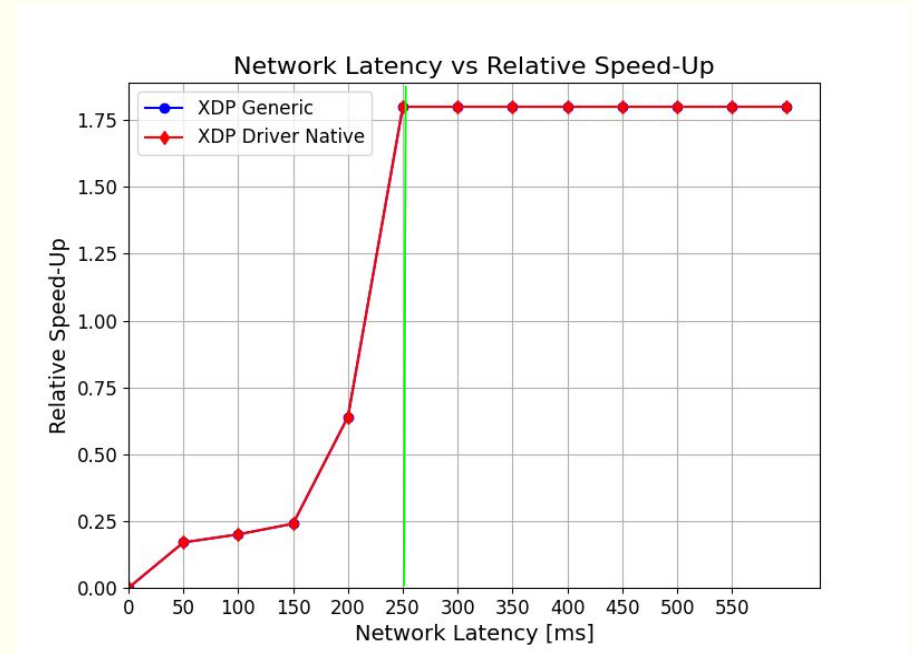



Results

- Optimal Speed-Up:

Acc. Network latency \geq Container initialization time

- Container “warm start” is the limit



Q&A