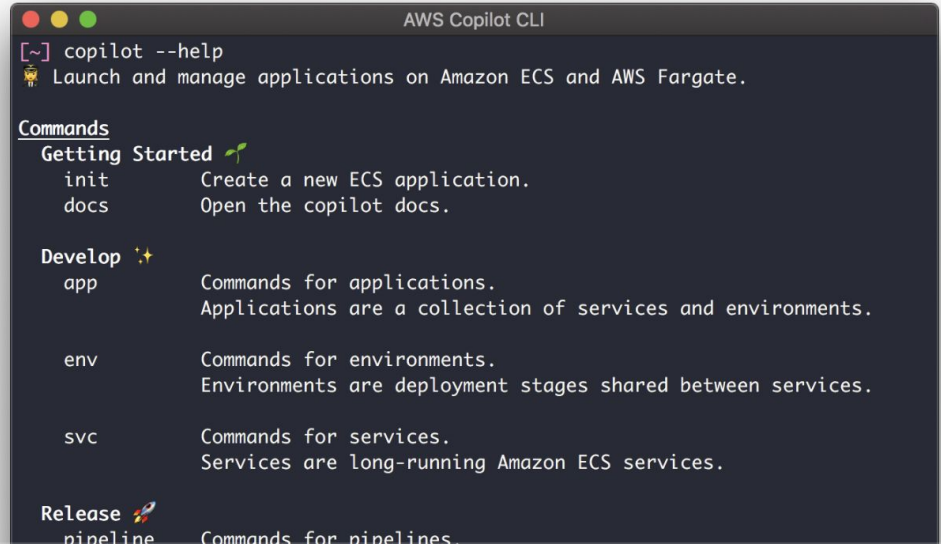


AWS Copilot

Primož Vampelj

What is AWS Copilot?

An AWS CLI tool that enables developers to deploy and manage simple applications backed by AWS Fargate or App Runner to AWS with supporting infrastructure



```
AWS Copilot CLI
[~] copilot --help
👤 Launch and manage applications on Amazon ECS and AWS Fargate.

Commands
Getting Started 🌱
  init      Create a new ECS application.
  docs      Open the copilot docs.

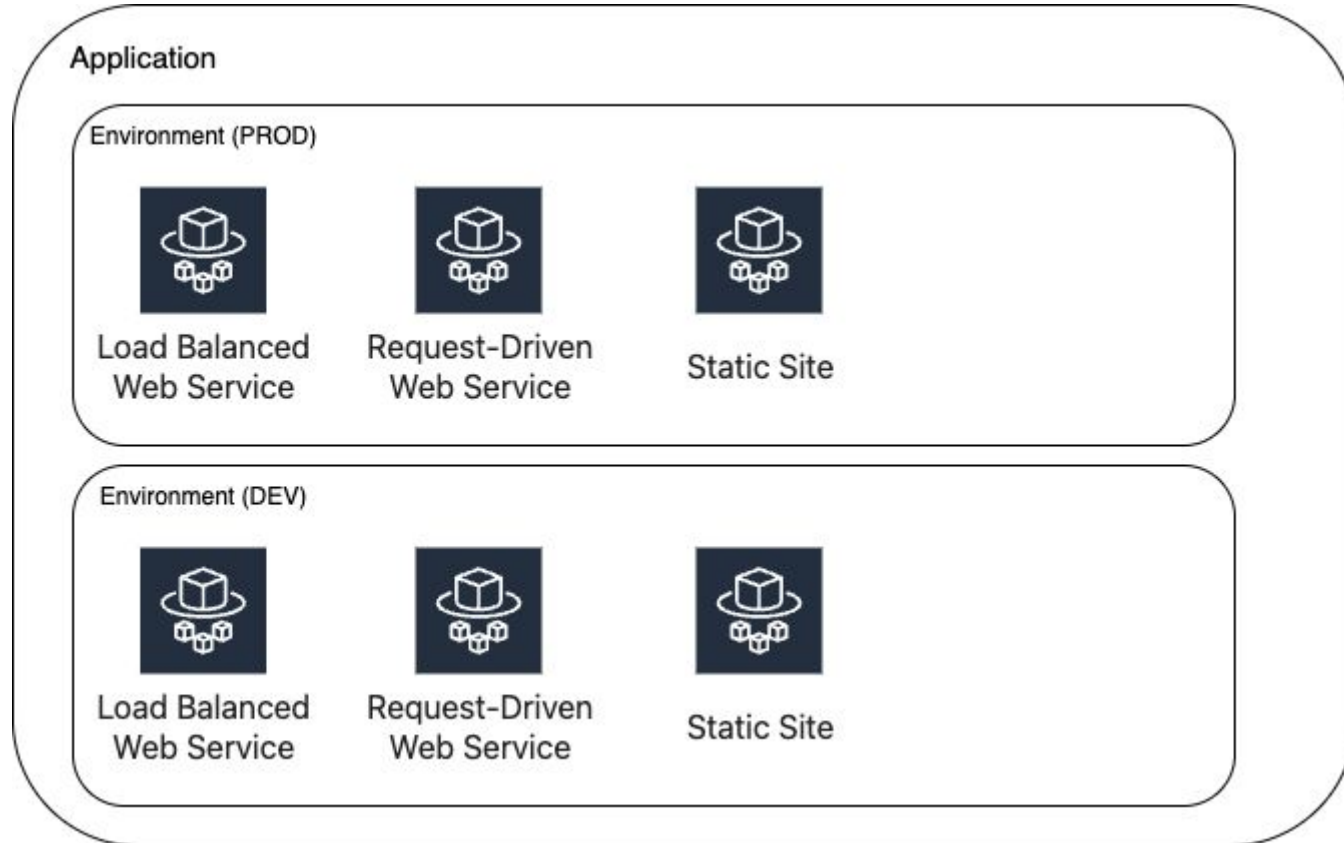
Develop 🌟
  app       Commands for applications.
            Applications are a collection of services and environments.

  env       Commands for environments.
            Environments are deployment stages shared between services.

  svc       Commands for services.
            Services are long-running Amazon ECS services.

Release 🚀
  pipeline  Commands for pipelines.
```

High level overview

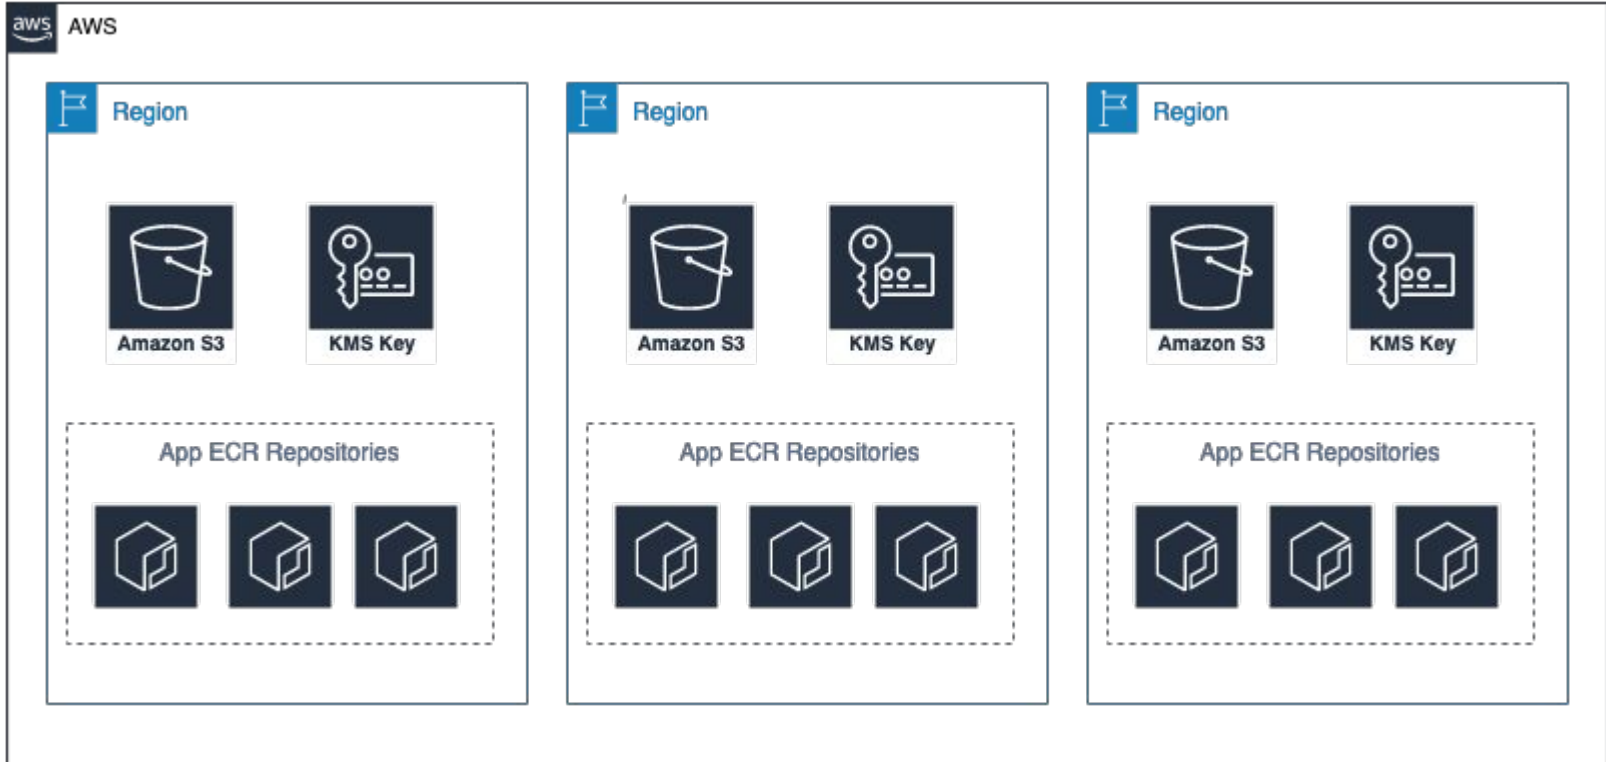


Application

- Groups connected services, environments and pipelines
- How to start?

```
$ copilot app init \
  --domain my-awesome-app.aws \
  --resource-tags department=MyDept,team=MyTeam \
  --permissions-boundary my-pb-policy
```

What gets created



```
$ copilot app show
```

About

Name	vote
Version	v1.1.0
URI	vote-app.aws

Environments

Name	AccountID	Region
----	-----	-----
test	0000000000000	us-east-1

Workloads

Name	Type	Environments
----	-----	-----
collector	Load Balanced Web Service	prod
aggregator	Backend Service	test, prod

Pipelines

Name

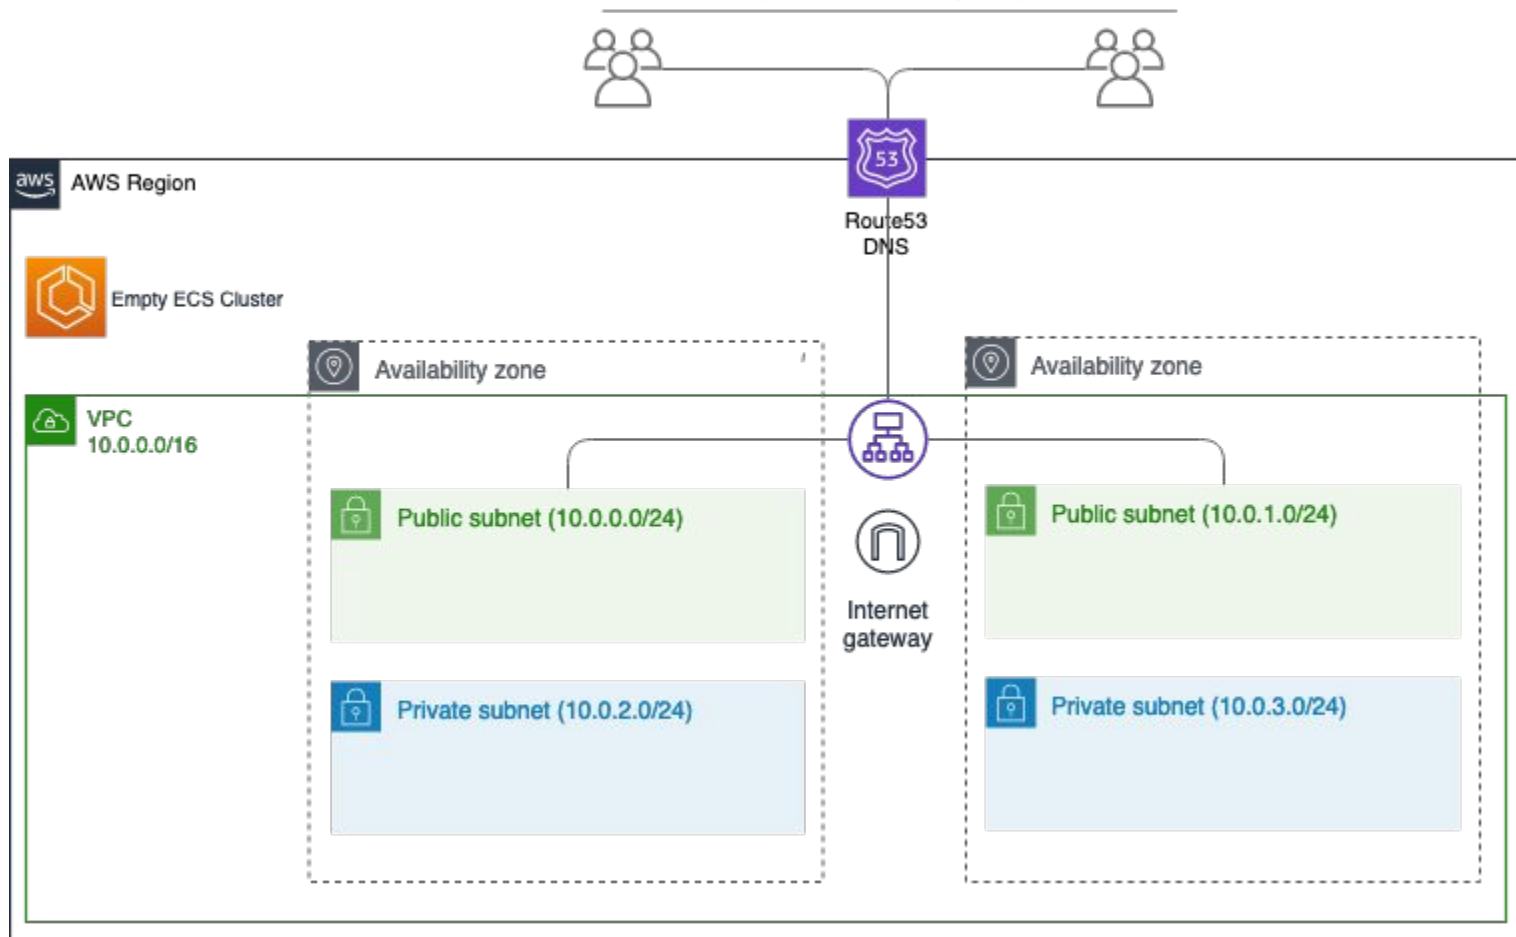
Environments

- You can create it from scratch or import your existing infrastructure
- Define your:
 - Networking (VPC, Subnets)
 - Load Balancer SG + certificates
- Environments have their own manifest file and can be re-deployed

```
> copilot env init
Environment name: test
Credential source: [profile vampelj_admin]

Would you like to use the default configuration for a new environment?
  - A new VPC with 2 AZs, 2 public subnets and 2 private subnets
  - A new ECS Cluster
  - New IAM Roles to manage services and jobs in your environment
[Use arrows to move, type to filter]
  Yes, use default.
  Yes, but I'd like configure the default resources (CIDR ranges, AZs).
> No, I'd like to import existing resources (VPC, subnets).
```

Environment Infrastructure



Services

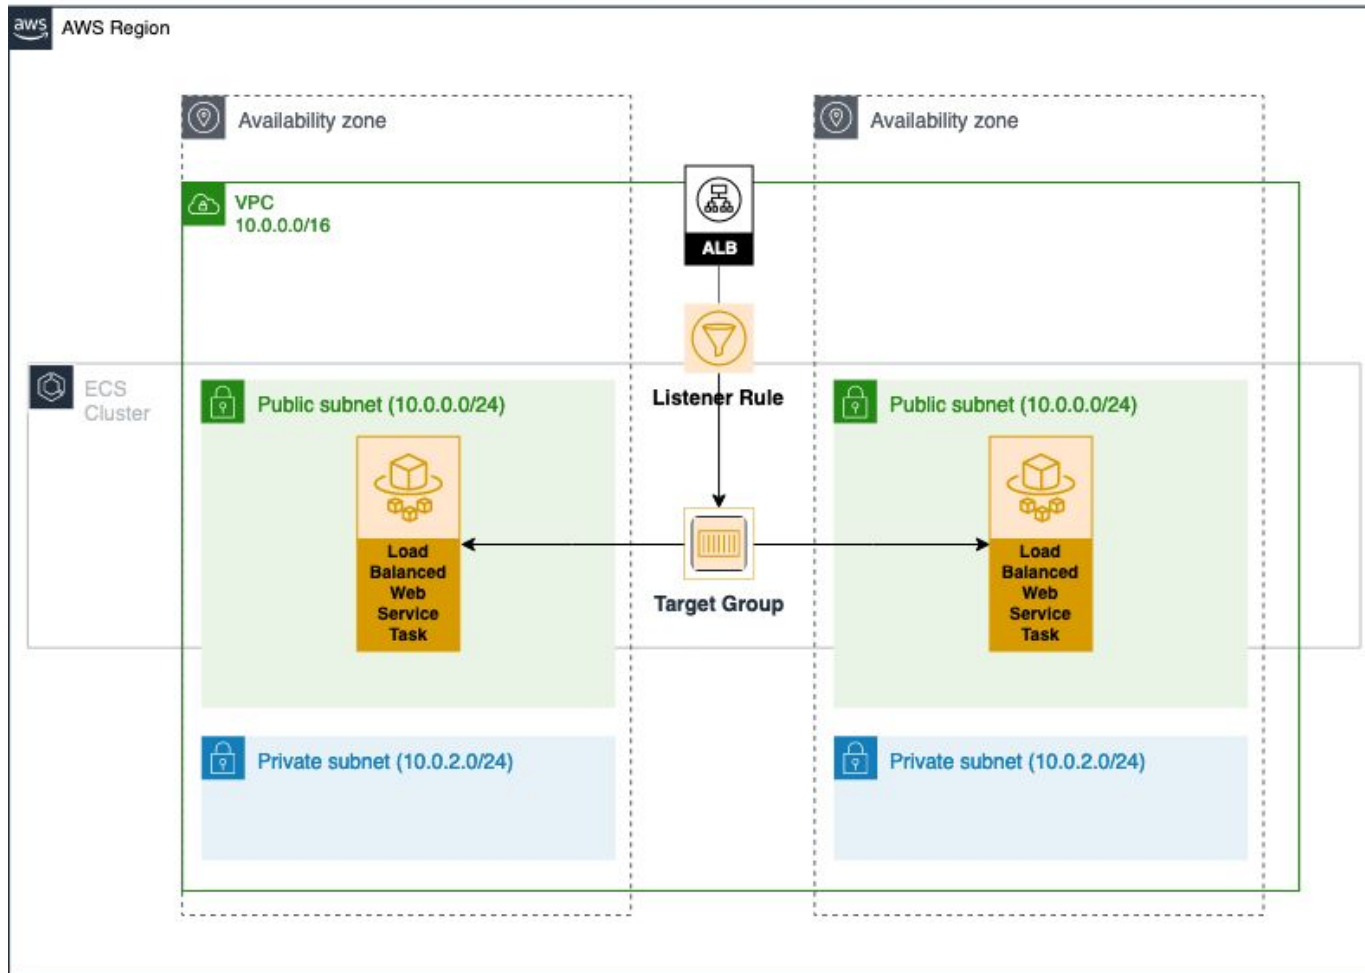
Internet facing:

- Request-Driven Web Service (App Runner)
- Static Site (S3 + CloudFront)
- **Load Balanced Web Service** (ALB and/or NLB, ECS Fargate)

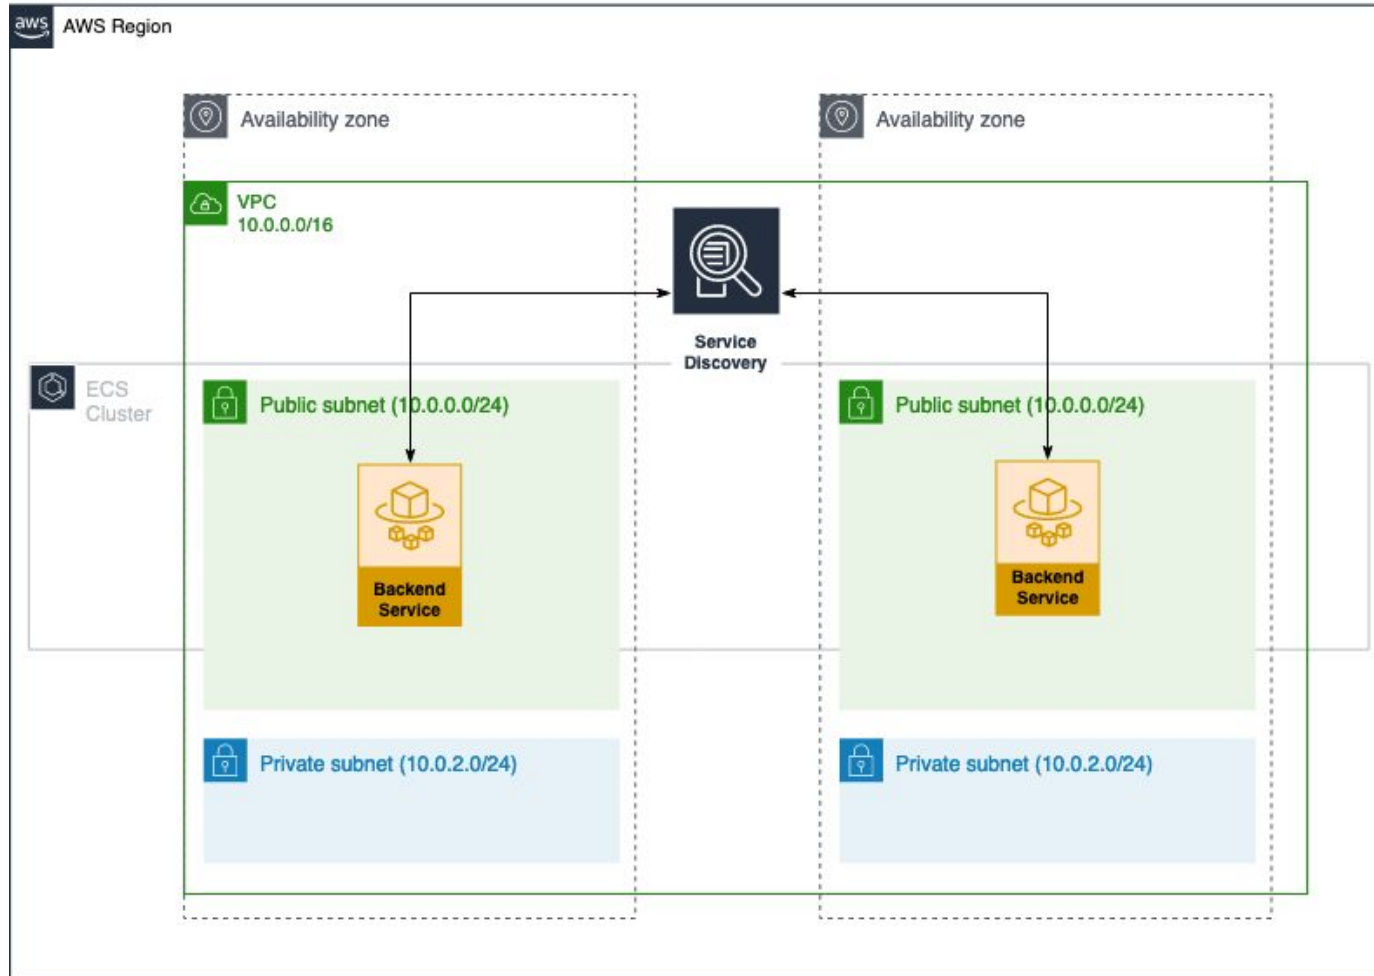
Private:

- Backend Service (ECS Fargate)
- Worker Service (SQS, ECS Fargate)

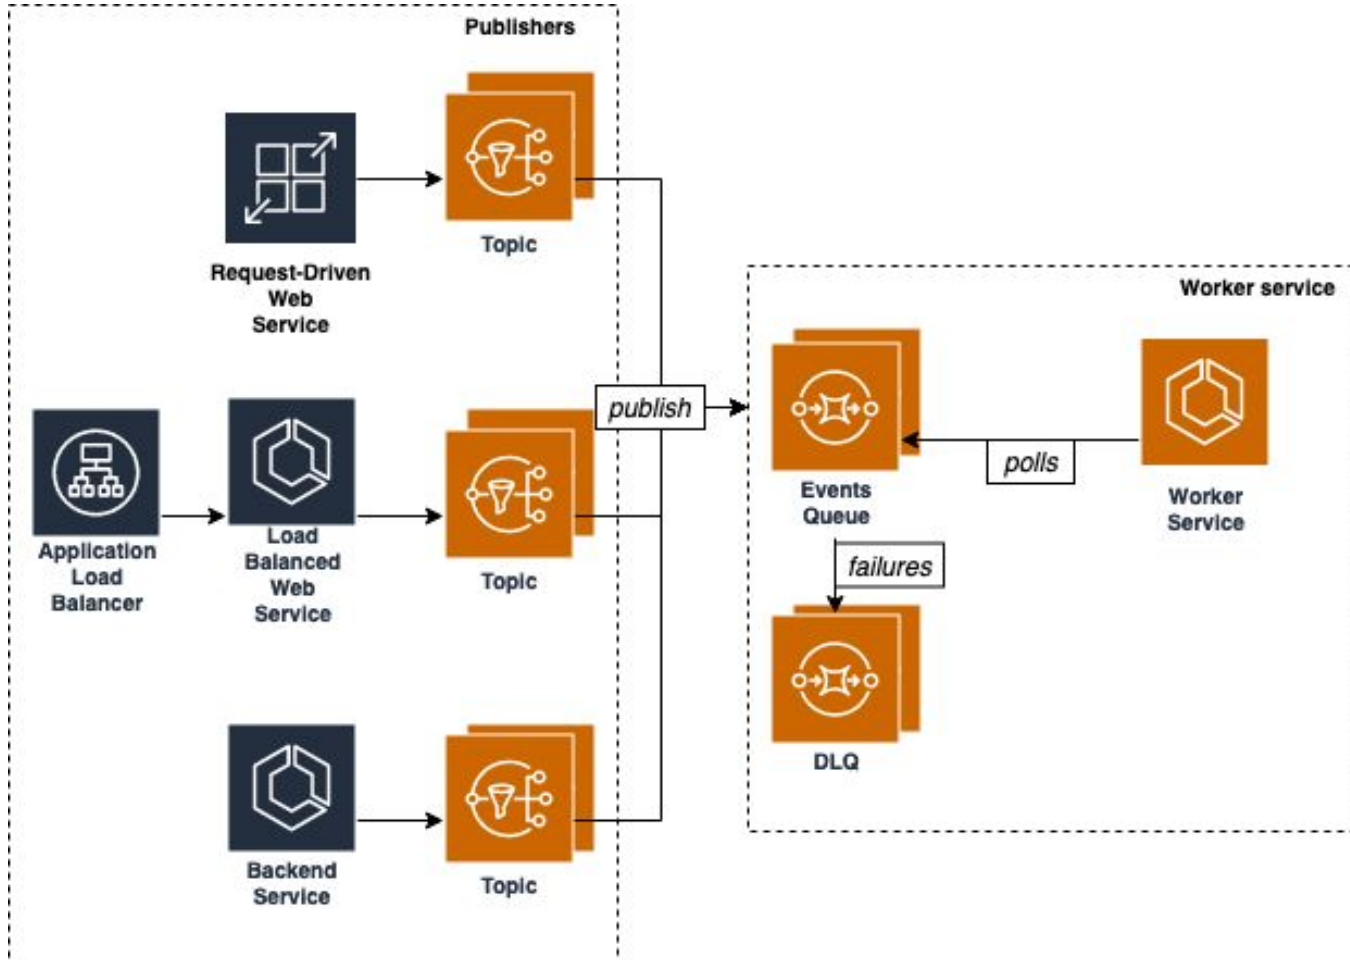
Load Balanced Web Service Infrastructure



Backend Service Infrastructure



Worker service





```
name: front-end
type: Load Balanced Web Service
```

```
image:
```

```
# Path to your service's Dockerfile.
```

```
build: ./Dockerfile
```

```
# Port exposed through your container to route traffic to it.
```

```
port: 8080
```

```
http:
```

```
# Requests to this path will be forwarded to your service.
```

```
# To match all requests you can use the "/" path.
```

```
path: '/'
```

```
# You can specify a custom health check path. The default is "/"
```

```
# healthcheck: '/'
```

```
# Number of CPU units for the task.
```

```
cpu: 256
```

```
# Amount of memory in MiB used by the task.
```

```
memory: 512
```

```
# Number of tasks that should be running in your service.
```

```
count: 1
```

```
# Optional fields for more advanced use-cases.
#
variables:                                # Pass environment variables as key value pairs.
  LOG_LEVEL: info

#secrets:                                # Pass secrets from AWS Systems Manager (SSM) Pa
#  GITHUB_TOKEN: GH_SECRET_TOKEN        # The key is the name of the environment variabl
#                                         # the value is the name of the SSM parameter.

# You can override any of the values defined above by environment.
environments:
  prod:
    count: 2                             # Number of tasks to run for the "prod" environment.
```

Pipelines

- AWS Code deploy pipeline
- Source => Build => Deploy (to multiple environments)

```
$ copilot pipeline init  
$ git add copilot/ && git commit -m "Adding pipeline artifacts" && git push  
$ copilot pipeline deploy
```

```
1  name: test-test-pipeline
2  version: 1
3
4  # This section defines your source, changes to which trigger your pipeline.
5  source:
6    provider: Bitbucket
7    properties:
8      branch: develop
9      repository: https://bitbucket.org/vampelj/test-test
10 stages:
11   - name: dev
12     pre_deployments:
13       db_migration:
14         buildspec: copilot/pipelines/test-test/db_migrations.yml
15     post_deployments:
16       celery_beat:
17         buildspec: copilot/pipelines/test-test/celery_beat.yml
18   - name: prod
19     requires_approval: true
20     pre_deployments:
21       db_migration:
22         buildspec: copilot/pipelines/test-test/db_migrations.yml
23     post_deployments:
24       celery_beat:
25         buildspec: copilot/pipelines/test-test/celery_beat.yml
26
```


Extras

- Storage: S3, EFS, DynamoDB, Aurora
- Secrets: SSM Parameter store / Secret Manager
- Support for sidecars (logging, Grafana Alloy)
- Overrides (YAML Patch Overrides / CDK Overrides)
- Jobs (scheduled)

The good, the bad and the ugly

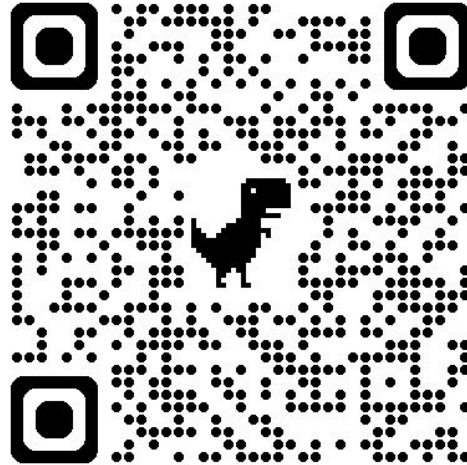
- + Easy to use and fast to deploy
- + Empowers developers
- + Great for simpler apps / infrastructure
- o Can become cumbersome to manage a lot of manifest files
- o Cloudformation takes its time to deploy
- Failed deploys can be a pain to fix

Tips

- Watch out for version mismatch
- READ THE DOCUMENTATION !!!
- If it seems wrong, it probably is!

Want to learn more

<https://aws.github.io/copilot-cli/>



Questions?

Thank you !